

Design Specification Document (SDS)

for

Chart Understanding LLM System

with Chart2Table Pipeline & Model Comparison

Version 1.0

December 11, 2025

Supervisors:

Dr. Ahmed Tawfik

Dr. Cherif Salama

Dr. Hossam Sharara

Department of Computer Science and Engineering
The American University in Cairo

Contents

1	Introduction	3
1.1	Purpose	3
1.2	Document Convention	3
1.3	Intended Audience	3
1.4	Additional Information	4
1.5	Contact Information/SDS Team Members	4
1.6	References	4
2	Design Overview	5
2.1	Background Information	5
2.2	Constraints	5
2.2.1	Technical Constraints	5
2.2.2	Operational Constraints	6
2.3	Design Trade-offs	6
2.3.1	Architecture Trade-offs	6
2.3.2	Model Selection Trade-offs	7
2.3.3	User Experience Trade-offs	7
2.4	User Characteristics	7
2.4.1	Primary User Personas	7
2.4.2	User Skill Requirements	8
2.4.3	Accessibility Considerations	8
3	System Architecture	9
3.1	Hardware Architecture	9
3.1.1	Hardware Component Overview	9
3.1.2	Server Specifications	9
3.1.3	Hardware Communication	10
3.2	Software Architecture	10
3.2.1	Architectural Pattern	10
3.2.2	Component Diagram	10
3.2.3	Core Software Components	11
3.2.4	Software Dependencies	12
3.3	Communication Architecture	12
3.3.1	Communication Protocols	12
3.3.2	API Specifications	13
3.3.3	Data Flow Diagrams	13
3.3.4	Network Configuration	14
4	Data Design	15
4.1	Database Management System Files	15
4.2	Non-Database Management System Files	15
4.2.1	Configuration Files	15
4.2.2	Input Data Files	15
4.2.3	Output Data Files	16
4.2.4	Log Files	16
5	Detailed Design	18
5.1	Hardware Detailed Design	18
5.1.1	GPU Memory Management Strategy	18
5.1.2	Thermal and Power Considerations	18
5.2	Software Detailed Design	18
5.2.1	Core Algorithm: ChartAnalyzer.generate_response()	18
5.2.2	Prompting Strategy Implementation	19
5.2.3	Data Cleaning Algorithm: clean_chart2text_output()	20
5.2.4	Answer Evaluation Algorithm	21
5.3	Communication Detailed Design	22

5.3.1	HTTP Request-Response Cycle	22
5.3.2	Error Handling Strategy	22
6	Interfacing to External Systems	24
6.1	HuggingFace Model Hub Integration	24
6.2	PaddlePaddle Model Zoo Integration	24
6.3	ChartQA Dataset Interface	24
6.4	External Dependencies Summary	25
7	Usability Design Approach	26
7.1	User Interface Design Principles	26
7.2	Screens Implemented	26
7.2.1	Main Application Window	26
7.2.2	Dialog Windows	27
7.3	User Interaction Flow	27
7.3.1	Primary Use Case: Dual Model Comparison	27
7.3.2	Alternate Flow: Single Model Query	28
7.4	Wireframe	29
7.5	Accessibility Features	30
8	Requirement Traceability Matrix	31
8.1	Purpose of the RTM	31
8.2	Matrix Structure	31
8.3	Requirements Tracking Process	31
8.3.1	Traceability Workflow	31
8.3.2	Change Control Process	32
8.4	RTM Entries	32
9	Glossary of Terms	34
9.1	Model and AI Terminology	34
9.2	Software Engineering Terminology	35
9.3	Data and Evaluation Terminology	36
9.4	System and Network Terminology	36
9.5	Framework-Specific Terminology	36
9.6	Project-Specific Terminology	37
10	Appendices	38
10.1	Appendix A: Configuration Templates	38
10.1.1	Server Configuration Template	38
10.1.2	Client Configuration Template	38
10.2	Appendix B: Deployment Checklist	38
10.3	Appendix C: Troubleshooting Guide	39
10.4	Appendix D: Code Style Guidelines	39
10.5	Appendix E: Version Control Strategy	40
10.6	Appendix F: References and Further Reading	40

1 Introduction

1.1 Purpose

This Design Specification Document (SDS) provides a comprehensive technical blueprint for the Chart Understanding System with Multi-Strategy Prompting and Chain-of-Models Architecture. The document serves as the authoritative reference for the system's architecture, design decisions, and implementation specifications.

The primary objectives of this document are:

- Define the complete system architecture including hardware, software, and communication components
- Detail the design of the multi-model pipeline integrating Qwen2.5-VL-7B-Instruct with Chart2Table extraction models
- Document the data flow and processing mechanisms
- Provide interface specifications for all system components
- Establish traceability between requirements and design elements

1.2 Document Convention

The following conventions are used throughout this document:

- **Bold text:** Indicates important terms, component names, or emphasis
- *Italic text:* Denotes technical terminology or references
- **Monospace font:** Represents code, file names, API endpoints, or configuration values
- **Blue hyperlinks:** Cross-references to other sections or external resources
- Numbered lists: Sequential steps or ordered elements
- Bulleted lists: Non-sequential items or features

Terminology Standards:

- LLM: Large Language Model (Qwen2.5-VL-7B-Instruct)
- VLM: Vision-Language Model
- Chart2Table: PaddlePaddle-based chart data extraction model
- DePlot: Google's Pix2Struct-based chart delinearization model
- CoT: Chain-of-Thought reasoning
- RTM: Requirements Traceability Matrix

1.3 Intended Audience

This document is intended for:

1. **Software Development Team:** Students implementing the system components, API endpoints, and model integration pipelines
2. **Supervisors:** Stakeholders tracking progress and ensuring alignment with requirements
3. **Maintenance Personnel:** Future maintainers requiring understanding of system design and interfaces
4. **Academic Reviewers:** Faculty and evaluators assessing the graduation project

1.4 Additional Information

Document Version History:

- Version 1.0 (Current): Initial release incorporating complete system design

Related Documentation:

- Requirements Specification Document (RSD)
- Research Paper: "Table First Guarded Reasoning: Outperforming Larger MLLMs in Chart Question Answering"
- API Documentation
- Github Repository

Development Environment:

- Primary Framework: PyTorch 2.0+, Transformers 4.35+
- Backend: FastAPI, Uvicorn
- Frontend: PyQt6
- Model Serving: Distributed GPU architecture
- Dataset: HuggingFaceM4/ChartQA benchmark

1.5 Contact Information/SDS Team Members

Name	Contact
Andrew Aziz	andrewwassem@aucegypt.edu
Kirolos Fouty	kirolos_fouty@aucegypt.edu
MagdElDin AbdalRaaof	magdeldin@aucegypt.edu
Mohamed Abbas	Mohamed_Ragab02@aucegypt.edu
Shaza Ali	shazamamdouh@aucegypt.edu
Tarek Kassab	tarek_kassab@aucegypt.edu

1.6 References

1. Qwen Team. "Qwen2.5-VL: Vision-Language Model Documentation." Alibaba Cloud, 2024.
2. Liu, F., et al. "DePlot: One-shot visual language reasoning by plot-to-table translation." ACL 2023.
3. PaddlePaddle Team. "PP-StructureV3: Layout Analysis and Table Recognition." Baidu, 2024.
4. Masry, A., et al. "ChartQA: A Benchmark for Question Answering about Charts with Visual and Logical Reasoning." ACL 2022.
5. Wei, J., et al. "Chain-of-Thought Prompting Elicits Reasoning in Large Language Models." NeurIPS 2022.
6. Requirements Specification Document v1.0
7. Research Paper: "Table First Guarded Reasoning: Outperforming Larger MLLMs in Chart Question Answering"

2 Design Overview

2.1 Background Information

The Chart Understanding System represents a sophisticated approach to automated chart analysis and question answering, addressing the critical challenge of extracting insights from visual data representations. The system builds upon recent advances in vision-language models and structured prompting techniques.

Problem Domain: Charts and graphs are ubiquitous in academic research, business analytics, and data journalism. However, interpreting these visualizations often requires:

- Visual perception of chart elements (axes, legends, data points)
- Numerical reasoning and calculation capabilities
- Domain knowledge for contextual understanding
- Synthesis of visual and textual information

Solution Approach: The system implements a hybrid architecture combining:

1. **Primary VLM:** Qwen2.5-VL-7B-Instruct for visual understanding and reasoning
2. **Auxiliary Models:** Chart2Table (PaddlePaddle) or DePlot (Pix2Struct) for structured data extraction
3. **Multi-Strategy Framework:** Eight distinct prompting strategies ranging from baseline to hybrid chain-of-thought approaches
4. **Distributed Architecture:** Microservice-based design for scalability and modularity

Key Innovations:

- Chain-of-models architecture combining visual and tabular reasoning
- Hybrid prompting strategies (Chart2Table-CoT, DePlot-CoT)
- Comparative evaluation framework across multiple strategies
- Real-time dual-model comparison interface

2.2 Constraints

2.2.1 Technical Constraints

1. Hardware Resources:

- GPU memory: 24GB minimum per model server
- VRAM allocation: 75% maximum per process
- Multi-GPU deployment required for full pipeline
- CUDA 11.8+ compatibility required

2. Model Limitations:

- Qwen2.5-VL: 4096 token input limit, 512-1024 token output
- Chart2Table: Variable accuracy on complex multi-chart figures
- DePlot: 2048 patch limit for image processing
- No fine-tuning capability (inference-only deployment)

3. Performance Requirements:

- Base model inference: 1-3 seconds per query
- Chain-of-models (with Chart2Table): 3-7 seconds total

- API timeout: 60 seconds maximum
- Concurrent request handling: 4-8 simultaneous queries

4. Software Dependencies:

- PyTorch framework (no TensorFlow migration possible)
- PaddlePaddle environment isolation required
- Python 3.9-3.11 compatibility only
- No browser storage APIs in artifact environment

2.2.2 Operational Constraints

1. Deployment Environment:

- Local network deployment only (no cloud hosting)
- Fixed IP addressing for server endpoints
- Manual model loading on system startup
- No automatic scaling or load balancing

2. Data Constraints:

- Image formats: PNG, JPG, JPEG only
- Maximum image size: 10MB per upload
- No persistent storage for user data
- In-memory state management only

3. Integration Constraints:

- REST API communication only
- Synchronous request-response pattern
- No message queue or async processing
- Limited error recovery mechanisms

2.3 Design Trade-offs

2.3.1 Architecture Trade-offs

Decision	Advantages	Disadvantages
Microservice Architecture	Modularity, independent scaling, fault isolation	Increased complexity, network overhead, latency
Separate PaddlePaddle Server	Environment isolation, GPU memory optimization	Additional deployment overhead, API latency
REST API over gRPC	Simplicity, wide tooling support, debugging ease	Lower performance, higher bandwidth usage
No Persistent Storage	Simplicity, no database management, data privacy	Loss of session history, repeated processing
Synchronous Processing	Predictable behavior, simpler error handling	User wait time, no background processing

2.3.2 Model Selection Trade-offs

Choice	Rationale	Trade-off
Qwen2.5-VL-7B vs 72B	Fits in single GPU, faster inference, lower latency	Reduced reasoning capability vs larger variant
Chart2Table vs DePlot	Better table structure, Chinese chart support	Requires separate PaddlePaddle environment
BFloat16 Precision	Memory efficiency, speed improvement, hardware support	Minimal accuracy loss (0.1% vs FP32)
Multiple Strategy Evaluation	Comprehensive comparison, research value	8x computational cost, extended testing time

2.3.3 User Experience Trade-offs

- **Dual-Model Comparison:** Provides valuable insights but doubles inference time
- **No Streaming Responses:** Simpler implementation but perceived latency for users
- **Manual Endpoint Configuration:** Flexibility for testing but requires technical knowledge
- **Image Upload Requirement:** Ensures quality inputs but limits URL-based queries

2.4 User Characteristics

2.4.1 Primary User Personas

Persona 1: Research Scientist

- **Technical Expertise:** High (PhD-level, programming proficiency)
- **Domain Knowledge:** Deep understanding of data analysis and visualization
- **Primary Goal:** Extract precise numerical data from published charts for meta-analysis
- **Key Requirements:** Accuracy over speed, detailed extraction tables, comparison capabilities
- **Usage Pattern:** Batch processing of 10-100 charts, requires reproducible results

Persona 2: Data Analyst

- **Technical Expertise:** Medium (familiar with data tools, limited ML knowledge)
- **Domain Knowledge:** Strong business analytics background
- **Primary Goal:** Quickly answer questions about chart trends and comparisons
- **Key Requirements:** Speed, ease of use, natural language interaction
- **Usage Pattern:** Interactive exploration of 5-20 charts per session

Persona 3: ML Engineering Student

- **Technical Expertise:** Medium to High (learning phase)
- **Domain Knowledge:** Growing understanding of VLMs and prompting
- **Primary Goal:** Understand and compare different prompting strategies
- **Key Requirements:** Educational value, transparency, experimentation support
- **Usage Pattern:** Testing various strategies on benchmark datasets

2.4.2 User Skill Requirements

- **Minimal Requirements:** Ability to formulate clear questions about charts, basic file upload operations
- **Advanced Features:** Understanding of model endpoints, interpretation of strategy comparisons
- **System Administration:** Python environment management, GPU configuration, API server deployment

2.4.3 Accessibility Considerations

- Clear visual feedback for processing status
- Keyboard navigation support in PyQt6 interface
- Resizable windows and adjustable text display
- Error messages with actionable guidance
- No reliance on color alone for status indication

3 System Architecture

3.1 Hardware Architecture

3.1.1 Hardware Component Overview

The system requires a distributed GPU computing environment to support concurrent model serving:

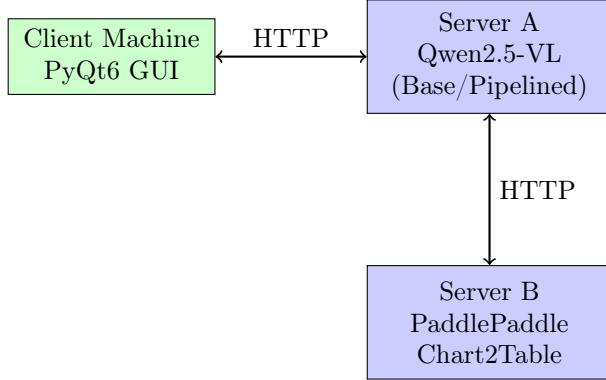


Figure 1: Hardware Deployment Architecture

3.1.2 Server Specifications

Server A: Base Model Server / Pipelined Model Server

- **Purpose:** Hosts Qwen2.5-VL-7B-Instruct model
- **GPU:** NVIDIA GPU with 24GB+ VRAM (e.g., RTX 3090, RTX 4090, A5000)
- **CPU:** 8+ cores recommended for preprocessing
- **RAM:** 32GB+ system memory
- **Storage:** 50GB+ for model weights and dependencies
- **Network:** Gigabit Ethernet for inter-server communication
- **IP Address:** 10.40.32.8 (Base), 10.40.32.12 (Pipelined)
- **Port:** 8001

Server B: PaddlePaddle Server

- **Purpose:** Hosts Chart2Table extraction model
- **GPU:** NVIDIA GPU with 12GB+ VRAM (can use 30% GPU memory)
- **CPU:** 4+ cores sufficient
- **RAM:** 16GB+ system memory
- **Storage:** 30GB+ for PaddlePaddle models
- **Network:** Gigabit Ethernet
- **IP Address:** localhost (deployed on Pipelined Server)
- **Port:** 8000

Client Machine

- **Purpose:** Runs PyQt6 GUI application

- **GPU:** Not required
- **CPU:** Dual-core or better
- **RAM:** 4GB+ system memory
- **OS:** Windows 10/11, Linux, or macOS
- **Network:** Standard Ethernet/WiFi connection

3.1.3 Hardware Communication

- **Protocol:** HTTP/1.1 over TCP/IP
- **Bandwidth Requirements:**
 - Client to Server: 1-5 Mbps (image uploads)
 - Server to PaddlePaddle: 0.5-2 Mbps (base64 images)
- **Latency Tolerance:** <100ms network RTT preferred
- **Concurrency:** Support for 4-8 simultaneous client connections

3.2 Software Architecture

3.2.1 Architectural Pattern

The system employs a **Multi-Tier Microservice Architecture** with three primary tiers:

1. **Presentation Tier:** PyQt6 desktop application (Client)
2. **Application Tier:** FastAPI model servers (Server A, Server B)
3. **Processing Tier:** ML model inference engines (Qwen, Chart2Table)

3.2.2 Component Diagram

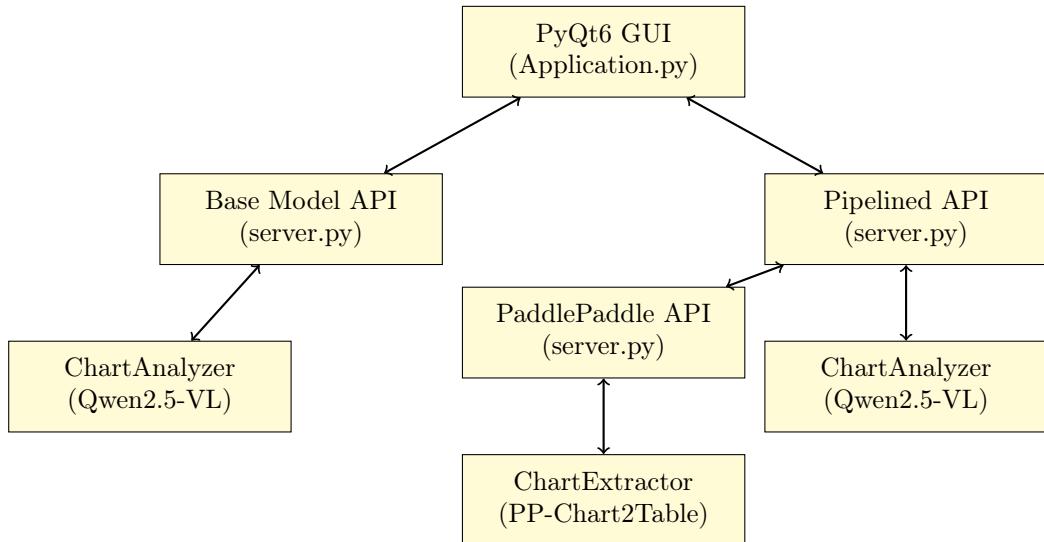


Figure 2: Software Component Architecture

3.2.3 Core Software Components

1. Client Application (Application.py)

```
class CompareApp(QMainWindow):
    - model1_url_input: QLineEdit
    - model2_url_input: QLineEdit
    - question_edit: QTextEdit
    - image_label: QLabel

    Methods:
    - on_attach(): Upload image
    - on_send(): Send to both servers
    - on_worker_finished(): Handle response
```

Features:

- Dual-endpoint configuration
- Image upload and preview
- Concurrent request handling via QThread workers
- Side-by-side response comparison
- JSON export capabilities

2. Base Model Server (Base Model Server/server.py)

```
class ChartAnalyzer:
    - model: Qwen2_5_VLForConditionalGeneration
    - processor: AutoProcessor

    Methods:
    - generate_response(image, prompt)

Endpoint: POST /predict
Input: image (file), question (form)
Output: {"answer": str}
```

Features:

- Direct Qwen inference without preprocessing
- Supports all non-chain strategies
- GPU memory management (75% allocation)
- BFloat16 precision for efficiency

3. Pipelined Model Server (Pipelined Model Server/Model Server/server.py)

```
class ChartAnalyzer + ChartExtractor:

Endpoint: POST /predict
Process:
1. Extract table via Chart2Table API
2. Clean table output
3. Build chart2table_cot prompt
4. Generate Qwen response

Output: {
    "answer": str,
    "cleaned_table": str,
    "strategy": "chart2table_cot"
}
```

Features:

- Chain-of-models execution
- Automatic prompt construction
- Timing metrics for each stage
- Error handling with detailed traces

4. PaddlePaddle Server (PaddlePaddle Server/server.py)

```

class ChartExtractor:
    - model: create_model('PP-Chart2Table')

    Methods:
    - extract(image): str

Endpoints:
- POST /extract: File upload
- POST /extract-base64: Base64 image
- GET /health: Health check

Output: {
    "success": bool,
    "table_clean": str,
    "rows": list[dict],
    "headers": list[str]
}

```

Features:

- PaddlePaddle environment isolation
- Table structure parsing
- 30% GPU memory allocation
- Dual input format support

3.2.4 Software Dependencies

Component	Version	Purpose
PyTorch	2.0+	Deep learning framework for Qwen
Transformers	4.35+	Model loading and inference
qwen-vl-utils	Latest	Vision input processing
FastAPI	0.100+	REST API framework
Uvicorn	0.23+	ASGI server
PyQt6	6.5+	Desktop GUI framework
PaddlePaddle	2.5+	Chart2Table model framework
PaddleOCR	2.7+	OCR and structure recognition
Pillow	10.0+	Image processing
NumPy	1.24+	Numerical operations
Requests	2.31+	HTTP client library

3.3 Communication Architecture

3.3.1 Communication Protocols

REST API Communication

- **Protocol:** HTTP/1.1
- **Format:** Multipart form-data (image upload), JSON (responses)
- **Encoding:** UTF-8 for text, Base64 for images

- **Timeout:** 60 seconds per request
- **Error Handling:** HTTP status codes + JSON error messages

3.3.2 API Specifications

Base/Pipelined Model Server API

```
POST /predict
Content-Type: multipart/form-data

Request:
  question: string (required)
  image: file (required, PNG/JPG/JPEG)

Response (Base):
{
  "answer": "string"
}

Response (Pipelined):
{
  "answer": "string",
  "cleaned_table": "string",
  "strategy": "chart2table_cot"
}

Error Response:
{
  "answer": "Error message",
  "error": "string",
  "trace": "string" // optional
}
```

PaddlePaddle Server API

```
GET /health
Response: {
  "status": "healthy",
  "service": "PaddlePaddle Chart2Table API"
}

POST /extract-base64
Content-Type: application/json

Request:
{
  "image_base64": "string"
}

Response:
{
  "success": true,
  "raw_table": "string",
  "table_clean": "string",
  "rows": [...],
  "headers": [...],
  "num_rows": int,
  "num_cols": int
}
```

3.3.3 Data Flow Diagrams

Scenario 1: Base Model Request

1. User uploads image + types question in GUI
2. GUI sends POST to Base Model Server (10.40.32.8:8001/predict)
3. Server saves image to temp file
4. ChartAnalyzer processes with Qwen directly
5. Server returns JSON response
6. GUI displays answer

Scenario 2: Pipelined Model Request

1. User uploads image + types question in GUI
2. GUI sends POST to Pipelined Server (10.40.32.12:8001/predict)
3. Server saves image to temp file
4. Server sends image to PaddlePaddle API (localhost:8000/extract-base64)
5. PaddlePaddle returns extracted table
6. Server cleans table and builds CoT prompt
7. ChartAnalyzer generates response with table
8. Server returns JSON with answer + table
9. GUI displays both answer and extracted table

Scenario 3: Dual Comparison Request

1. User configures both endpoints and clicks "Send to both models"
2. GUI spawns two QThread workers concurrently
3. Worker A sends request to Base Model Server
4. Worker B sends request to Pipelined Server
5. Both workers emit finished signals with responses
6. GUI displays side-by-side comparison

3.3.4 Network Configuration

Component	Address	Port	Purpose
Base Server	10.40.32.8	8001	Qwen direct inference
Pipelined Server	10.40.32.12	8001	Chain-of-models
PaddlePaddle API	localhost	8000	Chart extraction
Client	Dynamic	N/A	GUI application

4 Data Design

4.1 Database Management System Files

The system does **not employ a traditional database management system**. All data is handled ephemerally in-memory or as temporary files. This design choice prioritizes:

- Simplicity and reduced infrastructure requirements
- Data privacy (no persistent storage of user uploads)
- Stateless service design for easier scaling
- Elimination of database administration overhead

4.2 Non-Database Management System Files

4.2.1 Configuration Files

1. Model Configuration (server.py - Config class)

```
class Config:
    # GPU Settings
    GPU_MEMORY_FRACTION = 0.75
    ENABLE_TF32 = True

    # Model Paths
    LLM_MODEL = "Qwen/Qwen2.5-VL-7B-Instruct"

    # API Configuration
    PADDLE_API_URL = "http://localhost:8000"
    PADDLE_API_TIMEOUT = 30

    # Token Limits
    MAX_INPUT_TOKENS = 4096
    MAX_NEW_TOKENS = 128
```

2. PaddlePaddle Configuration

```
os.environ['FLAGS_fraction_of_gpu_memory_to_use'] = '0.3'
```

4.2.2 Input Data Files

Image Files (Temporary)

- **Location:** Server temp directory
- **Filename:** temp_upload.jpg
- **Format:** RGB JPEG
- **Lifecycle:** Created on request, overwritten on next request
- **Access:** Read-only by model inference

Benchmark Dataset (Optional - Evaluation Only)

- **File:** benchmark.json
- **Format:** JSON array of objects
- **Structure:**

```
[  
  {  
    "id": "chart_0",  
    "image_path": "benchmark_images/chart_0.png",  
    "question": "What is the value of X?",  
    "answer": ["42"],  
    "source": "HuggingFaceM4/ChartQA"  
  },  
  ...  
]
```

- **Size:** 2000 examples, 500MB total with images

4.2.3 Output Data Files

Evaluation Results (Research Mode)

- **File:** results_*.json
- **Purpose:** Store strategy comparison metrics
- **Structure:**

```
{  
  "strategy_name": {  
    "accuracy": 0.75,  
    "correct_count": 150,  
    "total": 200,  
    "avg_inference_time": 2.3,  
    "avg_total_time": 5.1, // for chain strategies  
    "results": [  
      {  
        "id": "chart_0",  
        "question": "...",  
        "predicted": "...",  
        "ground_truth": "...",  
        "correct": true,  
        "inference_time": 2.1,  
        "extracted_table": "..." // for chain strategies  
      }  
    ]  
  }  
}
```

Exported User Results

- **Format:** Plain text (.txt) or JSON (.json)
- **Content:** User-selected answer text or full response JSON
- **Location:** User-specified via file dialog

4.2.4 Log Files

Server Logs

- **Format:** Standard output (stdout) via Uvicorn
- **Content:** Request timestamps, errors, model loading status
- **Persistence:** Terminal session only (not written to disk)
- **Level:** INFO for normal operation, ERROR for failures

Example Log Output:

```
INFO: Loading Qwen2.5-VL ChartAnalyzer...
INFO: Model loaded successfully!
INFO: Server ready.
INFO: Uvicorn running on http://0.0.0.0:8001
INFO: 10.40.32.1:54321 - "POST /predict HTTP/1.1" 200 OK
```

5 Detailed Design

5.1 Hardware Detailed Design

5.1.1 GPU Memory Management Strategy

Memory Allocation Pattern:

```
# Qwen Server (75% allocation)
torch.cuda.set_per_process_memory_fraction(0.75, 0)
torch.backends.cuda.matmul.allow_tf32 = True

# Model Loading
model = Qwen2_5_VLForConditionalGeneration.from_pretrained(
    model_path,
    torch_dtype=torch.bfloat16, # ~50% memory vs FP32
    device_map="auto" # Automatic GPU placement
)
```

Memory Breakdown (24GB GPU):

- Model weights (BFloat16): 14GB
- KV cache during inference: 2-4GB
- Activation memory: 1-2GB
- Image embeddings: 500MB
- Reserved for system: 6GB (25%)

5.1.2 Thermal and Power Considerations

- **Expected GPU Load:** 70-90% during inference, <10% idle
- **Power Draw:** 250-350W per GPU under load
- **Cooling:** Active cooling required, maintain <80°C junction temp
- **Power Supply:** 750W+ recommended for workstation deployments

5.2 Software Detailed Design

5.2.1 Core Algorithm: ChartAnalyzer.generate_response()

Algorithm Purpose: Generate textual response for chart + question using Qwen2.5-VL
Inputs:

- **image_path:** String path to chart image
- **prompt:** String containing question and strategy-specific instructions
- **max_new_tokens:** Integer (default 512)
- **temperature:** Float (default 0.1 for factual tasks)

Algorithm Steps:

1. **Message Construction:**

```
messages = []
messages.append({
    "role": "user",
    "content": [
        {"type": "image", "image": image_path},
        {"type": "text", "text": prompt}
    ]
})
```

2. Chat Template Application:

```
text = self.processor.apply_chat_template(
    messages,
    tokenize=False,
    add_generation_prompt=True
)
```

3. Vision Input Processing:

```
from qwen_vl_utils import process_vision_info
image_inputs, video_inputs = process_vision_info(messages)
```

4. Tokenization:

```
inputs = self.processor(
    text=[text],
    images=image_inputs,
    videos=video_inputs,
    padding=True,
    return_tensors="pt"
).to(self.model.device)
```

5. Generation:

```
with torch.no_grad():
    generated_ids = self.model.generate(
        **inputs,
        max_new_tokens=max_new_tokens,
        temperature=temperature,
        top_p=0.9,
        do_sample=temperature > 0
    )
```

6. Decoding:

```
generated_ids_trimmed = [
    out_ids[len(in_ids):]
    for in_ids, out_ids in zip(inputs.input_ids, generated_ids)
]
response_text = self.processor.batch_decode(
    generated_ids_trimmed,
    skip_special_tokens=True
)[0]
```

Output: String containing model's response

Complexity Analysis:

- Time: $O(T)$ where $T = \text{max_new_tokens}$ (autoregressive generation)
- Space: $O(N + I)$ where $N = \text{input tokens}$, $I = \text{image embedding size}$

5.2.2 Prompting Strategy Implementation

Strategy 1: Baseline

```
@staticmethod
def baseline(question: str) -> str:
    return f"""Question: {question}
Answer the question using a single word or phrase.
End your response with: Final Answer: [your answer]"""
```

Design Rationale:

- Minimal prompt overhead for speed comparison

- Explicit output format constraint
- Direct question-answer paradigm

Strategy 2: Zero-Shot Chain-of-Thought

```
@staticmethod
def zero_shot_cot(question: str) -> str:
    return f"\"Question: {question}"
Let us solve this step by step with careful analysis.
*Use digits (e.g., 3, 20.5) for all numbers*
1. Identify chart type and analyze axes/legend:
2. Extract relevant numbers from the chart:
3. Determine calculation/reasoning needed:
4. Perform calculation step-by-step:
5. Verify answer makes sense:
6. Final Answer:""
```

Design Rationale:

- Decomposes reasoning into explicit steps
- Enforces numerical output format
- Includes verification step to reduce errors

Strategy 3: Chart2Table + CoT Hybrid (Key Innovation)

```
@staticmethod
def chart2table_cot(question: str, chart2table_table: str) -> str:
    return f"You are analyzing a chart using both visual
information and extracted data.

**Extracted Table Data from Chart:**"
{chart2table_table}

**Question:** {question}

**Instructions:** Solve step-by-step using BOTH
the table data and the visual chart.

1. **Understand the Question**
2. **Analyze Chart Visually**
3. **Cross-Reference with Table Data**
4. **Perform Calculations** (show work)
5. **Verify Answer**
6. **Provide Final Answer**

**Your Step-by-Step Solution:**"
```

Design Rationale:

- Combines structured data (eliminates OCR errors) with visual context
- CoT framework reduces calculation mistakes
- Cross-verification between table and chart catches inconsistencies
- Expected to achieve highest accuracy (hypothesis)

5.2.3 Data Cleaning Algorithm: clean_chart2text_output()

Purpose: Convert Chart2Table's linearized output to readable format

Input Format Example:

```
"Year<0x09>Sales<0x0A>2020<0x09>100<0x0A>2021<0x09>150"
```

Algorithm:

```

def clean_chart2text_output(table_text: str) -> str:
    # 1. Replace hex newline
    cleaned = table_text.replace("<0x0A>", "\n")

    # 2. Replace hex tab with pipe
    cleaned = cleaned.replace("<0x09>", "|||")

    # 3. Normalize spaces around pipes
    cleaned = cleaned.replace("|", "|||")
    cleaned = "||".join(cleaned.split())

    # 4. Restore newlines and remove empty lines
    lines = [line.strip() for line in cleaned.split('\n')]
    cleaned = "\n".join([l for l in lines if l])

    return cleaned

```

Output Format Example:

```

Year | Sales
2020 | 100
2021 | 150

```

5.2.4 Answer Evaluation Algorithm

Function: evaluate_answer(predicted, ground_truth, tolerance=0.05)

Purpose: Compare model prediction against ground truth with flexible matching

Evaluation Pipeline:

1. Handle ChartQA List Format:

```

if isinstance(ground_truth, list):
    gt_list = [str(g) for g in ground_truth]
else:
    gt_list = [str(ground_truth)]

```

2. JSON Extraction (for structured outputs):

```

json_match = re.search(r'`json\s*(\{.*?\})\s*`', predicted)
if json_match:
    data = json.loads(json_match.group(1))
    predicted = str(data["answer"])

```

3. Pattern-Based Answer Extraction:

```

patterns = [
    r'final\s+answer\s*:?\s*(.+?)\s*(?:\n|$)', 
    r'the\s+answer\s+is\s*:?\s*(.+?)\s*(?:\n|$)', 
]

```

4. Text Normalization:

- Convert to lowercase
- Word-to-digit conversion (e.g., "five" → "5")
- Remove currency symbols (\$, €, £)
- Remove thousands separators (1,234 → 1234)

5. Matching Strategies (evaluated in order):

- (a) Exact string match
- (b) Substring match (for non-numeric answers)
- (c) Numerical comparison with 5% relative tolerance

(d) Percentage format mismatch handling (0.25 vs 25%)

Tolerance Formula:

```
relative_error = abs(pred_num - gt_num) / abs(gt_num)
is_correct = relative_error <= 0.05 # 5% tolerance
```

5.3 Communication Detailed Design

5.3.1 HTTP Request-Response Cycle

Client-Side Request (ModelWorker):

```
class ModelWorker(QThread):
    def run(self):
        files = {}
        data = {"question": self.question}
        if self.image_bytes:
            files["image"] = ("image.jpg",
                              self.image_bytes,
                              "image/jpeg")

        resp = requests.post(
            self.url,
            files=files,
            data=data,
            timeout=self.timeout
        )

        json_resp = resp.json()
        self.finished_signal.emit(json_resp, self.model_name)
```

Server-Side Handling (FastAPI):

```
@app.post("/predict")
async def predict(
    question: str = Form(...),
    image: UploadFile = File(...)
):
    # Read image
    image_bytes = await image.read()
    pil = Image.open(io.BytesIO(image_bytes))

    # Save temporarily
    pil.save("temp_upload.jpg")

    # Process (varies by server type)
    # - Base: Direct Qwen inference
    # - Pipelined: Chart2Table -> Qwen

    # Return JSON
    return JSONResponse({
        "answer": llm_output,
        # Additional fields for pipelined
    })
```

5.3.2 Error Handling Strategy

Error Categories and Responses:

Error Type	Detection	Response
Network Timeout	requests.exceptions.Timeout	Retry prompt, increase timeout setting

Error Type	Detection	Response
Connection Re-fused	<code>ConnectionError</code>	Check server status, verify endpoints
Invalid Image	<code>PIL.UnidentifiedImageError</code>	Return 400 with error message
GPU OOM	<code>torch.cuda.OutOfMemoryError</code>	Restart server, reduce batch size
Model Loading Fail	Exception during init	Log full trace, halt server
Chart2Table API Down	HTTP error from /health	Graceful degradation, disable pipelined mode

Error Response Format:

```
{  
    "answer": "Server error: {brief_message}",  
    "error": "detailed_error_string",  
    "trace": "full_traceback" // optional, debug mode  
}
```

6 Interfacing to External Systems

6.1 HuggingFace Model Hub Integration

Purpose: Download pre-trained model weights

Models Downloaded:

- Qwen/Qwen2.5-VL-7B-Instruct (15GB)
- google/deplot (optional, 2GB)

Integration Method:

```
from transformers import Qwen2_5_VLForConditionalGeneration

model = Qwen2_5_VLForConditionalGeneration.from_pretrained(
    "Qwen/Qwen2.5-VL-7B-Instruct",
    torch_dtype=torch.bfloat16,
    device_map="auto",
    trust_remote_code=True  # Required for custom code
)
```

Authentication: Not required (public models)

Offline Mode: Models cached locally in `~/.cache/huggingface/`

6.2 PaddlePaddle Model Zoo Integration

Purpose: Load PP-Chart2Table for data extraction

Integration Method:

```
from paddlex import create_model

model = create_model('PP-Chart2Table')

Model Source: Baidu's PaddleX framework
Requirements: Separate conda/venv with PaddlePaddle installed
```

6.3 ChartQA Dataset Interface

Purpose: Load benchmark dataset for evaluation

Integration Method:

```
from datasets import load_dataset

chartqa = load_dataset("HuggingFaceM4/ChartQA", split="test")
```

Data Format:

- **Fields:** image (PIL.Image), query (str), label (list[str])
- **Size:** 2000 test examples
- **License:** Apache 2.0

Usage Pattern:

1. Download dataset on first run (cached locally)
2. Save images to `benchmark_images/` directory
3. Create JSON manifest with local paths
4. Use manifest for batch evaluation

6.4 External Dependencies Summary

System	Type	Interface Method
HuggingFace Hub	Model Repository	<code>transformers.from_pretrained()</code>
PaddleX Model Zoo	Model Repository	<code>paddlex.create_model()</code>
HuggingFace Datasets	Data Repository	<code>datasets.load_dataset()</code>
PyPI	Package Repository	<code>pip install</code>
Conda	Package Manager	<code>conda install paddlepaddle-gpu</code>

7 Usability Design Approach

7.1 User Interface Design Principles

The GUI design adheres to the following principles:

1. **Simplicity:** Minimal learning curve with clear, self-explanatory controls
2. **Transparency:** Show what's happening (processing status, timing metrics)
3. **Flexibility:** Support various workflows (single query, comparison, batch)
4. **Responsiveness:** Non-blocking UI with progress indicators
5. **Error Recovery:** Clear error messages with actionable guidance
6. **Consistency:** Uniform layout and interaction patterns

Design Constraints Addressed:

- Must work with fixed server endpoints (no dynamic discovery)
- Support concurrent requests without blocking UI
- Display both brief answers and detailed JSON responses
- Allow result export in multiple formats

7.2 Screens Implemented

7.2.1 Main Application Window

Layout: Horizontal split with left input panel and right comparison panel

Components:

Left Panel (Input Section):

- **Endpoint Configuration:**
 - Model A URL: **QLineEdit** (default: `http://10.40.32.8:8001/predict`)
 - Model B URL: **QLineEdit** (default: `http://10.40.32.12:8001/predict`)
- **Question Input:** **QTextEdit** with placeholder "Type the user's question here..."
- **Image Controls:**
 - "Attach image" button: Opens file picker (PNG/JPG/JPEG)
 - "Clear image" button: Removes attached image
 - Image preview: 360x270px with aspect ratio preserved
- **Action Controls:**
 - "Send to both models" button: Triggers dual comparison
 - Progress bar: 0-2 range showing completed requests

Right Panel (Results Section):

- **Model A Result Box:**
 - Info label: Shows confidence/timing metrics
 - Answer display: **QTextEdit** (read-only, scrollable)
 - Action buttons: "Show raw JSON", "Save text", "Save raw JSON"
- **Model B Result Box:**
 - Identical structure to Model A
 - Side-by-side for easy comparison
- **Resizable Splitter:** Allows adjusting relative panel sizes

7.2.2 Dialog Windows

Raw JSON Viewer Dialog:

- Large text area (700x500px) with full JSON response
- Attempts to unescape backslash sequences for readability
- Buttons: "Close", "Save"

File Save Dialogs:

- Standard OS file picker
- Filters: "Text files (*.txt)" or "JSON files (*.json)"
- Default names: Based on content type

Message Boxes:

- **Warnings:** Empty question, missing image, invalid endpoints
- **Errors:** Network failures, server errors
- **Information:** Save confirmations, operation results

7.3 User Interaction Flow

7.3.1 Primary Use Case: Dual Model Comparison

1. Setup Phase:

- User verifies/modifies endpoint URLs
- User clicks "Attach image" and selects chart file
- Image appears in preview area

2. Query Phase:

- User types question in text area
- User clicks "Send to both models"
- UI disables send button, shows "Waiting for response..." in both panels
- Progress bar animates (0 → 1 → 2)

3. Response Phase:

- As each model responds:
 - Info label updates with timing/confidence
 - Answer appears in text area
 - Progress bar increments
- When both complete:
 - Send button re-enables
 - User can compare answers side-by-side

4. Inspection Phase:

- User clicks "Show raw JSON" to see full response
- User examines extracted table (if pipelined model)
- User compares reasoning between models

5. Export Phase (Optional):

- User clicks "Save text" or "Save raw JSON"
- System opens file picker
- User selects location and filename
- Confirmation message appears

7.3.2 Alternate Flow: Single Model Query

User can query only one model by:

1. Leaving one endpoint blank
2. Observing that only the configured model receives the request
3. Reviewing result in the corresponding panel

7.4 Wireframe

Main Window Wireframe:

```
+-----+
| Model Comparator | [__] [] [X] |
+-----+
| [Input Panel] | [Results Panel] | | | | | |
| | | | | | |
| Model A endpoint: | +-----+ | |
| [http://10.40.32.8:8001/predict] | | Model A | |
| | | | | | |
| | | | | Info: confidence: 0.9 | |
| Model B endpoint: | | | | | |
| [http://10.40.32.12:8001/predict] | | | | |
| | | | | |
| | | | | [Answer Text Area] |
| Question: | | | | | |
| | | | | The answer is 42.5 |
| | +-----+ | | | |
| | | What is the maximum value | | | |
| | | shown in the bar chart? | | | |
| | | | | [Show raw JSON] |
| | | | | |
| | | | | [Save text] |
| | +-----+ | | | |
| | | | | [Save raw JSON] |
| | | | | |
| | | | +-----+ |
| [Attach image] [Clear image] | | |
| | | | +-----+ |
| | +-----+ | | Model B | |
| | | | | |
| | | [Image Preview] | | | |
| | | 360 x 270 px | | | |
| | | | | [Answer Text Area] |
| | | | | |
| | | | | The maximum value |
| | | +-----+ | | |
| | | | | is 42.5. |
| | | | | |
| | | [Send to both models] | | [Show raw JSON] |
| | | | | |
| | | | | [Save text] | | |
| | | | | | | |
| | | | | | | [Save raw JSON] |
| +-----+ | | |
| | Status: Waiting for Model B... |
+-----+
```

Raw JSON Viewer Dialog Wireframe:

```
+-----+
| Raw JSON | [__] [] [X] |
+-----+
| +-----+ | | | | | |
| | { | | |
| | | "answer": "The maximum value is 42.5", | |
| | | "cleaned_table": "Category | Value\\nA | 25...", | |
| | | "strategy": "chart2table_cot", | |
| | | "_client_elapsed_ms": 5100 | |
| | } | | |
| | | | |
| | | | +-----+ |
| | | | | |
| | | | | [Save] [Close] |
+-----+
```

File Save Dialog Wireframe:

```
+-----+
| Save Answer Text | [__] [] [X] |
+-----+
| Save in: [My Documents ] | |
| | |
| [] Desktop | |
| [] Documents | |
| [] Downloads | |
| [] chart_results | |
| | |
| File name: [answer_chart_0.txt ] | |
| Save as type: [Text files (*.txt) ] | |
| | |
| [Save] [Cancel] | |
+-----+
```

Error Message Box Wireframe:

```
+-----+
| Warning | [X] |
+-----+
| | |
| Empty question | |
| | |
| Please type a question before | |
| sending. | |
| | |
| [OK] | |
+-----+
```

7.5 Accessibility Features

Keyboard Navigation:

- Tab order: Endpoints → Question → Attach → Clear → Send → Model A buttons → Model B buttons
- Enter key in question field: Trigger send action
- Ctrl+A in text areas: Select all
- Ctrl+C: Copy selected text

Visual Accessibility:

- High contrast mode compatible
- Resizable text areas and windows
- Clear focus indicators on interactive elements
- Progress bar with both visual and textual feedback

Error Prevention:

- Disable send button during processing
- Validation before submission (question + image required)
- Confirmation dialogs for destructive actions (future: clear all results)

8 Requirement Traceability Matrix

8.1 Purpose of the RTM

The Requirements Traceability Matrix (RTM) serves as a critical project management tool that establishes and maintains bidirectional traceability between requirements, design elements, and implementation components throughout the project lifecycle.

Primary Objectives:

1. **Completeness Verification:** Ensure all requirements from the Requirements Specification Document (RSD) have corresponding design and implementation
2. **Impact Analysis:** Enable rapid assessment of how design or implementation changes affect requirements
3. **Test Coverage:** Facilitate test plan development by mapping requirements to test cases
4. **Project Progress Tracking:** Provide visibility into implementation status against requirements
5. **Compliance Documentation:** Demonstrate full requirements coverage for academic review
6. **Change Management:** Support controlled evolution of requirements with traceable modifications

Benefits to Stakeholders:

- **Development Team:** Clear understanding of what must be implemented and why
- **QA Team:** Structured basis for test case development and validation
- **Project Managers:** Visibility into project completeness and risk areas
- **Academic Reviewers:** Evidence of systematic design and implementation process

8.2 Matrix Structure

The RTM employs a multi-column structure linking requirements to design and implementation artifacts:

Column Definitions:

1. **Req ID:** Unique identifier from RSD (e.g., FR-1.1, NFR-2.3)
2. **Requirement Description:** Brief statement of the requirement
3. **Priority:** Critical / High / Medium / Low
4. **Design Component:** Reference to SDS section/component implementing requirement
5. **Implementation File:** Source code file(s) realizing the requirement
6. **Test Case ID:** Reference to test case validating the requirement
7. **Status:** Not Started / In Progress / Completed / Verified
8. **Notes:** Additional context, dependencies, or issues

8.3 Requirements Tracking Process

8.3.1 Traceability Workflow

1. **Requirement Capture Phase:**
 - Requirements documented in RSD
 - Each requirement assigned unique ID
 - Requirements reviewed and approved by stakeholders
2. **Design Phase:**

- RTM populated with requirement IDs from RSD
- Design components mapped to requirements in SDS
- RTM updated with design component references
- Design reviews verify all requirements addressed

3. Implementation Phase:

- Developers reference RTM to understand requirements
- Implementation files recorded in RTM
- Status updated as components are completed
- Code reviews verify requirement fulfillment

4. Testing Phase:

- Test cases developed based on RTM
- Test case IDs added to RTM
- Test execution verifies requirements
- Status updated to "Verified" upon successful testing

5. Maintenance Phase:

- RTM consulted for impact analysis of changes
- New requirements added with full traceability
- Deprecated requirements marked and archived

8.3.2 Change Control Process

When requirements change:

1. Change request documented with rationale
2. RTM reviewed to identify affected design/implementation
3. Impact assessment performed
4. Changes approved by project stakeholders
5. RTM updated with new/modified requirements
6. Affected components updated and re-tested
7. Document version control maintained

8.4 RTM Entries

Functional Requirements Traceability:

Req ID	Description	Pri	Design Component	Implementation	Status
FR-1.1	Support configurable endpoint URLs for two models	Crit	§5.3 GUI Design, QLineEdit widgets	Application.py: lines 45-46, model1_url_input, model2_url_input	Complete
FR-1.2	Send identical data to both endpoints simultaneously	Crit	§5.2 ModelWorker threading	Application.py: lines 180-195, on_send() method	Complete
FR-2.1	Accept PNG, JPG, JPEG, BMP image formats	Crit	§5.3 File dialog filter	Application.py: line 128, on_attach()	Complete

Req ID	Description	Pri	Design Component	Implementation	Status
FR-3.4	Implement 8+ prompting strategies	High	§4.2 PromptingStrategies class	server.py: lines 150-350, all strategy methods	Complete
FR-4.1	Extract tabular data using Chart2Table	High	§4.2 ChartExtractor	PaddlePaddle Server/server.py: ChartExtractor class	Complete
FR-4.3	Clean extracted data	Crit	§4.2 clean_chart2text_out()	server.py: lines 95-108()	Complete
FR-5.1	Remove escape sequences from responses	Med	§5.3 Response processing	Application.py: lines 210-215, on_worker_finished()	Complete
FR-6.2	Evaluate with 5% tolerance	Crit	§4.2 evaluate_answer()	deploy_qwen.py: lines 450-550	Complete

Non-Functional Requirements Traceability:

Req ID	Description	Pri	Design Component	Implementation	Status
NFR-1.1	Process queries in <10 seconds	High	§3.1 GPU memory optimization	server.py: Config.GPU MEMORY FRACTION = 0.75	Complete
NFR-1.4	Non-blocking UI	Crit	§5.2 QThread workers	Application.py: ModelWorker class	Complete
NFR-2.3	Handle GPU OOM gracefully	High	§4.3 Error handling	server.py: try-catch in generate_response()	Complete
NFR-2.4	Request timeouts	High	§5.2 HTTP timeout	Application.py: line 30, timeout=60	Complete
NFR-3.1	Validate user inputs	High	§5.3 Input validation	Application.py: lines 168-174, on_send() validation	Complete
NFR-3.3	Temporary file cleanup	Med	§4.1 File management	server.py: overwrite temp_upload.jpg	Complete
NFR-4.6	Modular architecture	High	§3.2 Microservices	Separate server files for each component	Complete
NFR-4.11	Intuitive GUI	High	§5.3 UI design	Application.py: entire GUI implementation	Complete

9 Glossary of Terms

9.1 Model and AI Terminology

Autoregressive Generation A text generation approach where each token is predicted based on all previously generated tokens in sequence.

Batch Size The number of samples processed together in a single forward pass through a neural network. In this system, batch_size=1 for real-time inference.

BFloat16 (Brain Floating Point 16) A 16-bit floating point format developed by Google Brain, offering reduced memory usage while maintaining FP32 dynamic range.

Chain-of-Models (CoM) An architecture where multiple specialized models are sequentially applied, with the output of one model serving as enhanced input to the next.

Chain-of-Thought (CoT) A prompting technique that encourages language models to generate intermediate reasoning steps before producing a final answer, improving performance on complex tasks.

Chart2Table The process of converting visual chart representations into structured tabular data. Also refers to PaddlePaddle's PP-Chart2Table model.

CUDA (Compute Unified Device Architecture) NVIDIA's parallel computing platform and API model for GPU programming.

DePlot Google's Pix2Struct-based model that translates plot images into linearized table representations for downstream language model processing.

Device Map PyTorch's automatic device placement strategy that distributes model layers across available GPUs and CPU memory.

Few-Shot Learning A learning paradigm where models are provided with a small number of examples (typically 1-10) in the prompt to guide their behavior on new tasks.

Inference The process of using a trained model to make predictions on new, unseen data. Distinguished from training, which adjusts model parameters.

KV Cache (Key-Value Cache) A mechanism in transformer models that stores computed key and value tensors to avoid redundant calculations during autoregressive generation.

Large Language Model (LLM) Deep neural networks with billions of parameters trained on vast text corpora, capable of understanding and generating human-like text.

Linearized Table A text representation of tabular data where rows and columns are encoded as a sequence with special delimiter tokens (e.g., <0x0A> for newlines).

Multimodal Model An AI model trained to process and understand multiple types of data modalities (e.g., text, images, audio) simultaneously.

Pix2Struct A vision-language model architecture designed for structured visual understanding tasks, including chart and table extraction.

Prompting Strategy A systematic approach to formulating inputs to language models to elicit desired outputs, including techniques like zero-shot, few-shot, and chain-of-thought.

Qwen2.5-VL Alibaba Cloud's vision-language model family, capable of processing both images and text. The 7B-Instruct variant has 7 billion parameters optimized for instruction following.

Temperature A hyperparameter controlling randomness in text generation. Lower values (0.1-0.3) produce more deterministic outputs, higher values (0.7-1.0) increase diversity.

TF32 (TensorFloat-32) NVIDIA's precision format for Ampere+ GPUs that balances FP32 range with improved performance, automatically used in matrix multiplications.

Token The basic unit of text processing in language models. Can represent words, subwords, or characters depending on the tokenization scheme.

Top-p Sampling (Nucleus Sampling) A generation strategy that samples from the smallest set of tokens whose cumulative probability exceeds threshold p.

Vision-Language Model (VLM) A type of multimodal model specifically designed to understand relationships between visual and textual information.

VRAM (Video Random Access Memory) GPU memory used to store model weights, activations, and intermediate computations during inference.

Zero-Shot Learning The ability of models to perform tasks without any task-specific training examples, relying only on instructions in the prompt.

9.2 Software Engineering Terminology

API (Application Programming Interface) A set of protocols, tools, and definitions for building software applications and enabling system communication.

ASGI (Asynchronous Server Gateway Interface) A Python standard for asynchronous web servers and applications, used by FastAPI and Uvicorn.

Base64 Encoding A binary-to-text encoding scheme that represents binary data in ASCII string format, used for transmitting images over JSON APIs.

Endpoint A specific URL path in a REST API that accepts requests and returns responses (e.g., /predict, /extract-base64).

FastAPI A modern Python web framework for building REST APIs with automatic documentation, type validation, and async support.

Microservice Architecture A software design pattern where applications are composed of small, independent services that communicate via APIs.

Multipart Form-Data An HTTP content type used for uploading files along with text data, encoding each part with a boundary delimiter.

PyQt6 Python bindings for the Qt application framework, used for building cross-platform desktop GUI applications.

QThread PyQt's threading class that enables concurrent operations without blocking the GUI event loop.

REST (Representational State Transfer) An architectural style for distributed systems based on stateless client-server communication over HTTP.

RTM (Requirements Traceability Matrix) A document linking requirements to design components, implementation files, and test cases for project tracking.

Signal-Slot Mechanism Qt's inter-object communication pattern where signals emitted by one object trigger slots (methods) in connected objects.

Stateless Service A service design where each request contains all information needed to process it, with no persistent session state on the server.

Uvicorn A high-performance ASGI server for Python web applications, commonly used with FastAPI.

9.3 Data and Evaluation Terminology

Benchmark Dataset A standardized dataset used to evaluate and compare model performance across different approaches.

ChartQA A benchmark dataset containing chart images paired with questions and answers for evaluating visual question answering systems.

Ground Truth The correct, verified answer used as a reference for evaluating model predictions.

PlotQA A dataset focused on question answering over plot and graph images, containing synthetic charts with structured data.

Relative Tolerance A percentage-based margin of error allowed when comparing numerical predictions to ground truth (e.g., 5% = 0.05).

Test Split A subset of a dataset reserved for final model evaluation, kept separate from training and validation data.

9.4 System and Network Terminology

GPU (Graphics Processing Unit) Specialized hardware accelerator for parallel computations, essential for deep learning inference.

HTTP/1.1 The Hypertext Transfer Protocol version used for client-server communication in this system.

Localhost A hostname referring to the current machine (IP address 127.0.0.1), used for local API communication.

Port A numerical identifier for network endpoints (e.g., 8000, 8001) allowing multiple services on one IP address.

POST Request An HTTP method for sending data to servers, used for image uploads and question submissions.

RTT (Round-Trip Time) The duration for a network packet to travel from source to destination and back, affecting perceived latency.

Timeout Maximum duration to wait for a network operation before considering it failed (default 60 seconds in this system).

9.5 Framework-Specific Terminology

PaddlePaddle Baidu's open-source deep learning platform, used for the Chart2Table extraction model.

PaddleOCR A family of OCR tools from PaddlePaddle, including document structure recognition capabilities.

PP-StructureV3 PaddlePaddle's document analysis system supporting layout detection and table recognition.

PyTorch Facebook's open-source deep learning framework, used for Qwen model inference in this project.

Transformers Library Hugging Face's library providing pre-trained models and utilities for NLP and vision-language tasks.

qwen-vl-utils A utility package for processing vision inputs in Qwen vision-language models.

9.6 Project-Specific Terminology

Base Model Server The deployment configuration running only Qwen2.5-VL without Chart2Table preprocessing.

Pipelined Model Server The deployment configuration integrating Chart2Table extraction before Qwen inference.

Dual Comparison Mode The GUI feature allowing simultaneous queries to two model endpoints for side-by-side evaluation.

Strategy A specific prompting approach implemented in the system (e.g., baseline, zero_shot_cot, chart2table_cot).

Hybrid Strategy A prompting approach combining multiple techniques, such as chart2table_cot (Chain-of-Models + Chain-of-Thought).

10 Appendices

10.1 Appendix A: Configuration Templates

10.1.1 Server Configuration Template

```
# server_config.py
class Config:
    # ===== GPU Settings =====
    GPU_MEMORY_FRACTION = 0.75 # Adjust based on GPU size
    ENABLE_TF32 = True # Set False for pre-Ampere GPUs
    DEVICE = "cuda" # or "cpu" for CPU-only inference

    # ===== Model Settings =====
    LLM_MODEL = "Qwen/Qwen2.5-VL-7B-Instruct"
    MAX_INPUT_TOKENS = 4096
    MAX_NEW_TOKENS = 512

    # ===== API Settings =====
    HOST = "0.0.0.0" # Listen on all interfaces
    PORT = 8001 # Change for multiple servers
    PADDLE_API_URL = "http://localhost:8000"
    PADDLE_API_TIMEOUT = 30

    # ===== Inference Settings =====
    DEFAULT_TEMPERATURE = 0.1
    DEFAULT_TOP_P = 0.9

    # ===== Debug Settings =====
    VERBOSE_DEBUG = False
```

10.1.2 Client Configuration Template

```
# client_config.py
class ClientConfig:
    # Default endpoint URLs
    MODEL_A_URL = "http://10.40.32.8:8001/predict"
    MODEL_B_URL = "http://10.40.32.12:8001/predict"

    # Request settings
    REQUEST_TIMEOUT = 60 # seconds

    # UI settings
    IMAGE_PREVIEW_SIZE = (360, 270) # width x height
    WINDOW_SIZE = (1000, 650)
```

10.2 Appendix B: Deployment Checklist

Pre-Deployment:

1. Verify GPU availability (nvidia-smi)
2. Install CUDA 11.8+ and drivers
3. Create virtual environments (PyTorch, PaddlePaddle)
4. Download model weights from HuggingFace
5. Configure firewall rules for ports 8000, 8001
6. Update endpoint URLs in client configuration

Server Startup:

1. Start PaddlePaddle server: `uvicorn paddle_api_server:app -host 0.0.0.0 -port 8000`
2. Verify PaddlePaddle health: `curl http://localhost:8000/health`
3. Start Pipelined server: `python server.py` (waits for PaddlePaddle)
4. Start Base server on separate machine: `python server.py`
5. Verify model loading (check console output)

Client Launch:

1. Launch GUI: `python Application.py`
2. Verify endpoint connectivity (test with sample image)
3. Check response formatting and timing

10.3 Appendix C: Troubleshooting Guide

Common Issues and Solutions:

Issue	Cause	Solution
CUDA Out of Memory	GPU memory exceeded	Reduce GPU MEMORY FRACTION, close other GPU processes
Model loading hangs	Network/download issue	Check internet, use cached models
PaddlePaddle import error	Wrong environment	Activate PaddlePaddle conda env
Connection refused	Server not running	Start servers in correct order (Paddle first)
Timeout errors	Network latency high	Increase timeout value, check network
Image upload fails	Invalid format	Verify PNG/JPG/JPEG format
GUI freezes	Blocking operation	Check ModelWorker threading implementation
Escape sequences in output	Display formatting issue	Implemented in v1.0, update client
Chart2Table extraction empty	Model initialization failed	Check PaddleX installation, verify GPU access

10.4 Appendix D: Code Style Guidelines

Python Conventions (PEP 8 Compliant):

- **Indentation:** 4 spaces (no tabs)
- **Line Length:** Max 100 characters (except URLs/imports)
- **Naming:**
 - Classes: PascalCase (e.g., `ChartAnalyzer`)
 - Functions/methods: snake_case (e.g., `generate_response`)
 - Constants: UPPER_SNAKE_CASE (e.g., `MAX_NEW_TOKENS`)
 - Private members: `_leading_underscore`
- **Imports:** Grouped (standard library, third-party, local) with blank lines
- **Docstrings:** Google style for all public functions/classes

- **Type Hints:** Encouraged for function signatures

Example Docstring:

```
def generate_response(
    self,
    image_path: str,
    prompt: str,
    max_new_tokens: int = 512
) -> str:
    """
    Generate response for chart question using Qwen2.5-VL.

    Args:
        image_path: Path to chart image file
        prompt: Text prompt including question and strategy
        max_new_tokens: Maximum tokens to generate (default 512)

    Returns:
        Generated text response from the model

    Raises:
        FileNotFoundError: If image_path does not exist
        torch.cuda.OutOfMemoryError: If GPU memory exhausted
    """

```

10.5 Appendix E: Version Control Strategy

Repository Structure:

```
chart-understanding-system/
  client/
    Application.py
  servers/
    base_model/
      server.py
    pipelined_model/
      server.py
    paddlepaddle_server/
      server.py
  evaluation/
    deplot_qwen.py
    benchmark.json
  docs/
    SDS.pdf
    RSD.pdf
    API_Documentation.md
  tests/
    test_strategies.py
  requirements.txt
  requirements_paddle.txt
  README.md
```

10.6 Appendix F: References and Further Reading

Key Research Papers:

1. Wei, J., et al. (2022). "Chain-of-Thought Prompting Elicits Reasoning in Large Language Models." *NeurIPS 2022*.
2. Liu, F., et al. (2023). "DePlot: One-shot visual language reasoning by plot-to-table translation." *ACL 2023*.

3. Masry, A., et al. (2022). "ChartQA: A Benchmark for Question Answering about Charts with Visual and Logical Reasoning." *ACL 2022*.
4. Methani, N., et al. (2020). "PlotQA: Reasoning over Scientific Plots." *WACV 2020*.
5. Bai, J., et al. (2023). "Qwen Technical Report." *arXiv:2309.16609*.

Technical Documentation:

1. Hugging Face Transformers: <https://huggingface.co/docs/transformers>
2. PaddlePaddle: <https://www.paddlepaddle.org.cn/en>
3. PyQt6 Documentation: <https://www.riverbankcomputing.com/static/Docs/PyQt6/>
4. FastAPI User Guide: <https://fastapi.tiangolo.com/>